

Building a Secure Embedded Kernel in Rust

Amit Levy^a, Branden Ghena^b, Bradford Campbell^b,
Pat Pannuto^b, Nicholas Matsakis^c, Prabal Dutta^b, Philip Levis^a

Platform Lab

May 8, 2017

^aStanford University ^bUniversity of Michigan ^cMozilla Research

The Internet of Things (IoT)



A Security Disaster

The Economist

World politics Business & finance Economics Science & technology Culture

Cyber-security

The internet of things (to be hacked)

Hooking up gadgets to the web promises huge benefits. But security must not be an afterthought

Jul 12th 2014 | From the print edition

Timekeeper

Like

217

Tweet

594

How the Internet of Things Could Kill You

By Fahmida Y. Rashid JULY 18, 2014 7:30 AM - Source: Tom's Guide US | 5 COMMENTS

Hacking the Fridge: Internet of Things Has Security Vulnerabilities

JESS SCANLON | MORE ARTICLES
JUNE 28, 2014

Philips Hue LED smart lights hacked, home blacked out by security researcher

By Sal Cangeloso on August 15, 2013 at 11:45 am | 7 Comments

- HP conducted a security analysis of IoT devices¹
 - ▶ 80% had privacy concerns
 - ▶ 80% had poor passwords
 - ▶ 70% lacked encryption
 - ▶ 60% had vulnerabilities in UI
 - ▶ 60% had insecure updates

¹http://fortifyprotect.com/HP_IoT_Research_Study.pdf

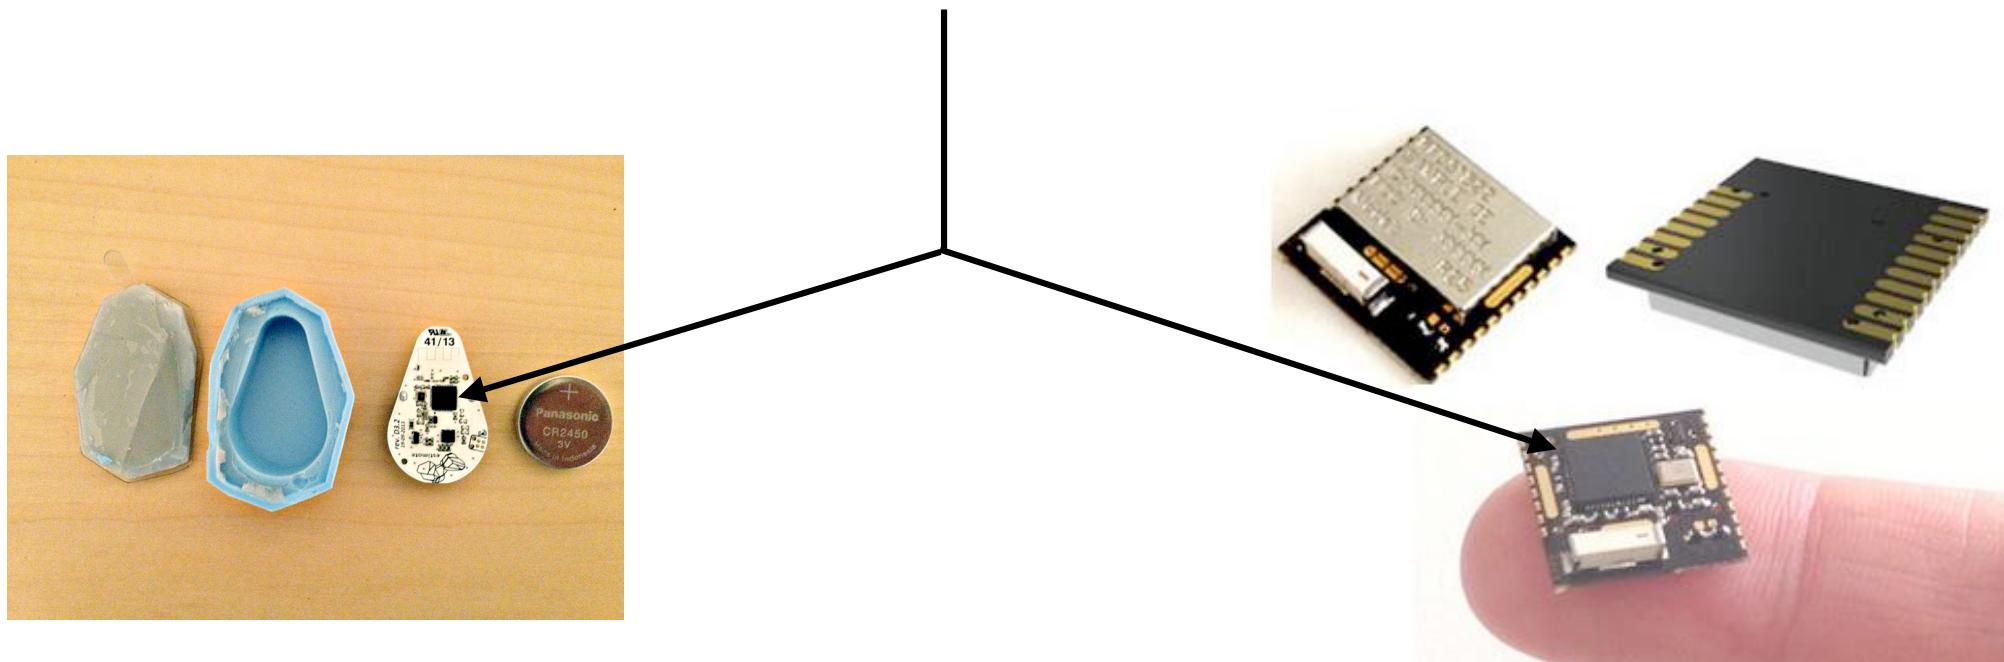
Let's design and build a secure
embedded kernel!

Outline

- Energy rules everything
- Kernel design goals
- Tock operating system design
 - ▶ Architecture: Rust, capsules, and processes
 - ▶ Rust challenges
 - ▶ Resource model: memory and CPU
- Some early observations

Embedded Devices

- Cortex M0+ with integrated 2.4GHz transceiver
 - ▶ Supports Bluetooth Low Energy
 - ▶ Two models: 32kB/256kB or 16kB/128kB
- DigiKey cost for 25,000: \$1.99



Energy Rules Everything

- Most things don't have wires connected to them
 - ▶ All energy comes from a battery
 - ▶ Energy harvesting is sometimes but often not possible
- For IoT devices, energy is the limiting resource
- Two costs dominate
 - ▶ Sleep/idle current
 - ▶ Radio communication

Sleep and Wake Up

- nRF51422
- CR2032: 225mAh @ 3.0V
 - ▶ Can drive system at 1.8V, use a buck power converter to step down voltage, in this slide assume we just run at 3.0V



| Mode | Current | Lifetime |
|-----------------------|---------|----------|
| Sleep | 10.2 µA | 920 days |
| CPU active | 4.4 mA | 2.1 days |
| Radio TX | 16 mA | 14 hours |
| Radio RX | 13.4 mA | 17 hours |
| 5Hz BLE advertisement | ~231 µA | 40 days |

Joules/Bit

- WiFi is more efficient than BLE and 802.15.4!
 - ▶ WiFi: 50Mbps (application), 600mW, 83MbpJ
 - ▶ 802.15.4: 200kbps (application), 60mW, 3.3MbpJ
 - ▶ BLE: 150kbps (application), 30mW, 5MbpJ
- But, transition times can dominate
 - ▶ Simpler radios wake up faster
- Simple case: send a 1kB packet
 - ▶ 1kB packet is 160μS in WiFi, 5ms in BLE
 - ▶ Suppose wakeup is 2ms WiFi, 200μS BLE
 - ▶ WiFi is awake 2.16ms (1.2mJ), BLE is awake 5.2ms (0.3mJ)
 - ▶ Costs less to send packet with BLE
 - ▶ 1kB is very large for IoT

Energy Rules Everything

- Moore's Law/Denard scaling don't drive designs

| | TelosB | imix |
|----------------------|---------------|-------------|
| Year | 2004 | 2017 |
| MCU | MSP430F161I | SAM4LC8CA |
| Sleep current | 0.2 μ A | 1.6 μ A |
| Word size | 16-bit | 32-bit |
| CPU clock | 8 MHz | 12-48 MHz |
| Flash | 48 kB | 512 kB |
| RAM | 10kB | 64kB |

13 years, 6-fold increase in RAM
and sleep current.

Outline

- Energy rules everything
- Kernel design goals
- Tock operating system design
 - ▶ Architecture: Rust, capsules, and processes
 - ▶ Rust challenges
 - ▶ Resource model: memory and CPU
- Some early observations

Dependability

- Embedded systems need to run for long periods of time without any intervention
- Place a high premium on *dependability*
- Often sacrifice other metrics, like speed or throughput, to improve dependability
 - ▶ E.g., if a FitBit crashes, reset it and lose all data; if that fails, return it to the manufacturer
 - ▶ What would you do if your USB authentication dongle was flaky/ crashed?

But the World is Changing...

“An embedded system is a computerized system that
is purpose-built for its application.”

Elicia White
Making Embedded Systems, O'Reilly

But the World is Changing...

“An embedded system is a computerized system that is purpose-built for its application.”

THE MUST-HAVE PEBBLE TIME APPS, WATCH FACES,
AND GAMES FOR YOUR WRIST

By Joshua Sherman — April 25, 2017 6:34 AM



Elicia White
Making Embedded Systems, O'Reilly

DON'T FALL
BEHIND

Stay current with a recap of
today's **Tech News** from
Digital Trends

Enter your email

SIGN UP

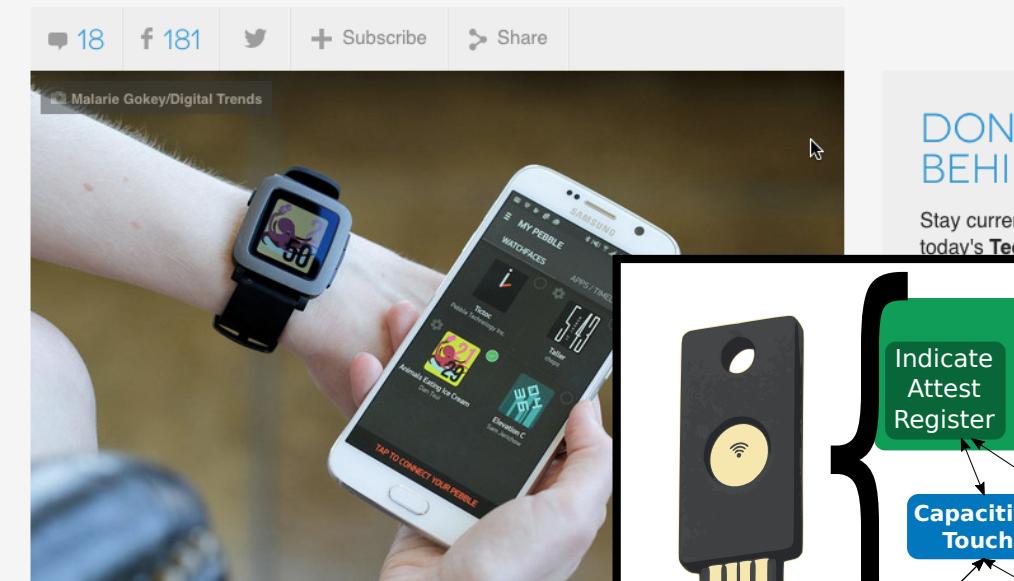
But the World is Changing...

“An embedded system is a computerized system that is purpose-built for its application.”

THE MUST-HAVE PEBBLE TIME APPS, WATCH FACES,
AND GAMES FOR YOUR WRIST

By Joshua Sherman — April 25, 2017 6:34 AM

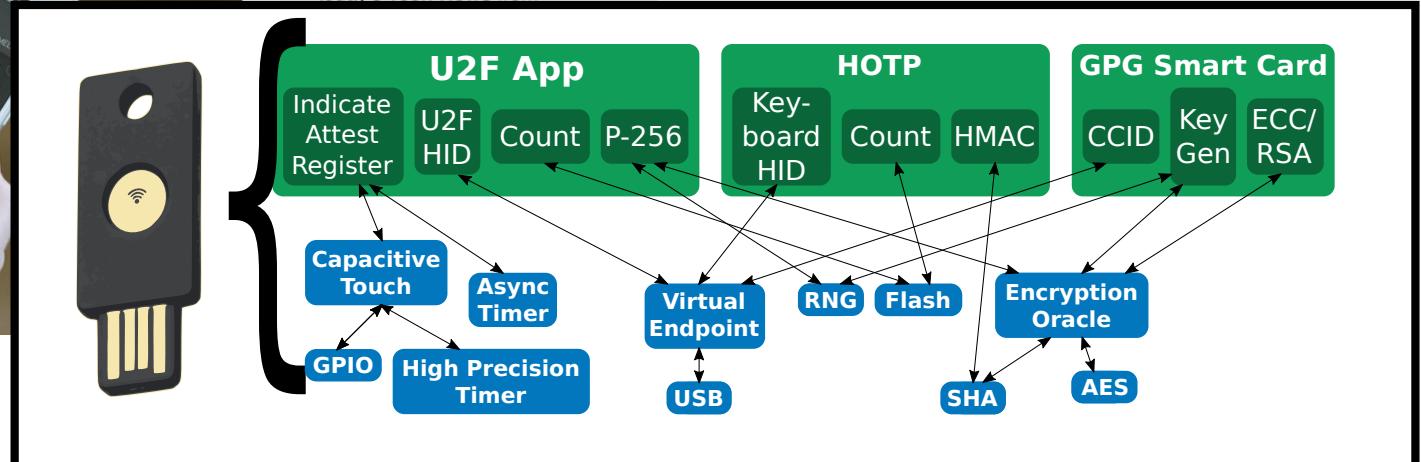
Elicia White
Making Embedded Systems, O'Reilly



18 f 181 Twitter + Subscribe Share

DON'T FALL BEHIND

Stay current with a recap of today's Tech News from



But the World is Changing...

“An embedded system is a computerized system that is purpose-built for its application.”

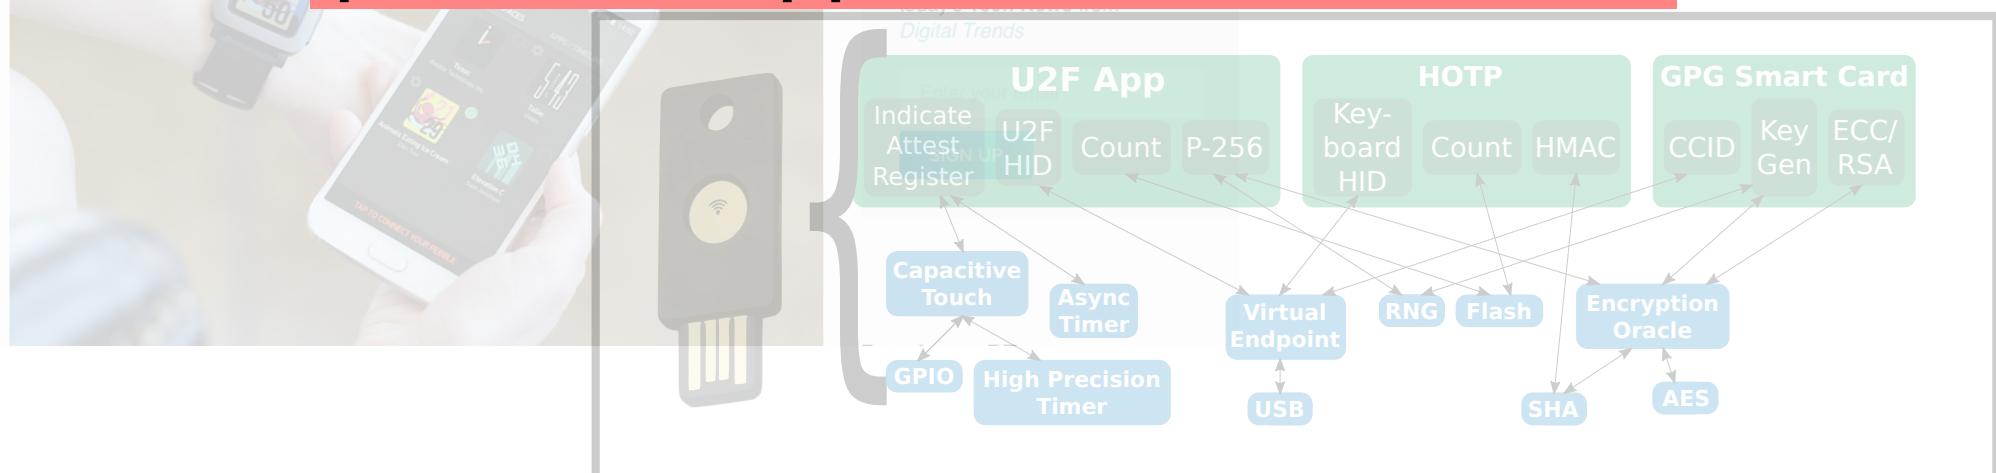
A new class of embedded devices, that act as platforms supporting loadable programs within a particular application domain.

THE MUST-HAVE
AND GAMES RANKED

By Joshua Sherman — April 25, 2017 6:34 AM

18 f 181

Mario Gokey/Digital Trends



5 Kernel Design Goals

- No Resource Exhaustion
 - ▶ Don't run out of memory at runtime
 - ▶ Application infinite loop doesn't hang system
- Memory Isolation
 - ▶ Protect independent pieces of code from each other
 - ▶ Historically impossible: written in C, no MMU
- Memory Efficiency
 - ▶ Conserve precious RAM
- Concurrency
 - ▶ Can perform multiple operations at once (energy efficiency)
- Loadable Applications
 - ▶ Can be dynamically reprogrammed without kernel reinstall

Previous Systems

| | No Resource Exhaustion | Memory Isolation | Memory Efficiency | Concurrency | Loadable Applications |
|-------------|------------------------|------------------|-------------------|-------------|-----------------------|
| Arduino | | | ✓ | | |
| TinyOS | | | ✓ | ✓ | |
| SOS | | | ✓ | ✓ | ✓ |
| TOSTThreads | ✓ | | | ✓ | ✓ |
| Contiki | | | ✓ | ✓ | ✓ |
| FreeRTOS | | | | ✓ | |
| RiotOS | | | ✓ | | |

Previous Systems

| | No Resource Exhaustion | Memory Isolation | Memory Efficiency | Concurrency | Loadable Applications |
|--|------------------------|------------------|-------------------|-------------|-----------------------|
| Arduino | | | ✓ | | |
| All of these are written in C, designed for 8-bit or 16-bit MCUs. Can more modern hardware and programming languages remove these tradeoffs? | | | | | |
| Contiki | | ✓ | ✓ | ✓ | |
| FreeRTOS | | | | ✓ | |
| RiotOS | | ✓ | | | |

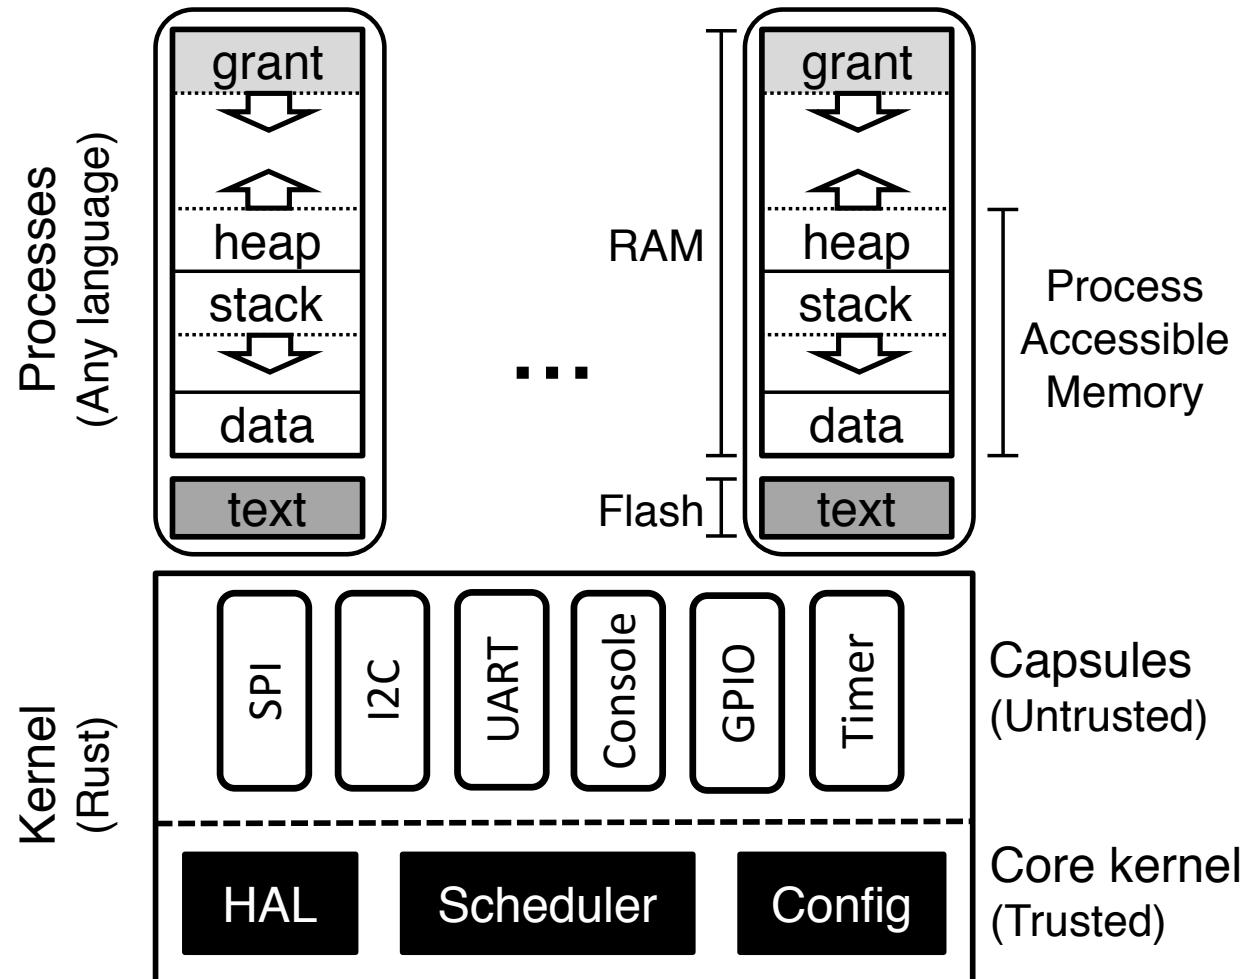
Outline

- Energy rules everything
- Kernel design goals
- Tock operating system design
 - ▶ Architecture: Rust, capsules, and processes
 - ▶ Rust challenges
 - ▶ Resource model: memory and CPU
- Some early observations

Tock Operating System

- Safe, multi-tasking operating system for memory-constrained devices
- Core kernel written in Rust, a safe systems language
 - ▶ Small amount of trusted code (can do unsafe things)
 - Rust bindings for memory-mapped I/O
 - Core scheduler, context switches
- Core kernel can be extended with *capsules*
 - ▶ Safe, written in Rust
 - ▶ Run inside kernel
- Processes can be written in any language (asm, C)
 - ▶ Leverage Cortex-M memory protection unit (MPU)
 - ▶ User-level, traps to kernel with system calls

Tock Architecture



Processes vs. Capsules

| Category | Capsule | Process |
|------------------------|----------------|----------------|
| Protection | Language | Hardware |
| Memory Overhead | None | Separate stack |
| Protection Granularity | Fine | Coarse |
| Concurrency | Cooperative | Preemptive |
| Update at Runtime | No | Yes |

Rust Safety

- Tackles two problems:
 - ▶ Thread safety (concurrent access)
 - ▶ Memory safety (address contains proper type)
- Rule 1: a memory location can have one read/write pointer or multiple read-only pointers
 - ▶ mutable references and references in Rust parlance
- Rule 2: a reference can only point to memory that is assured to outlive the reference
 - ▶ prevents dangling pointers

Rust Rule I

- Rule I: a memory location can have one read/write pointer or multiple read-only pointers
 - ▶ mutable references and references in Rust parlance

```
let mut x = 5;  
let y = &x;  
let z = &x;
```

OK

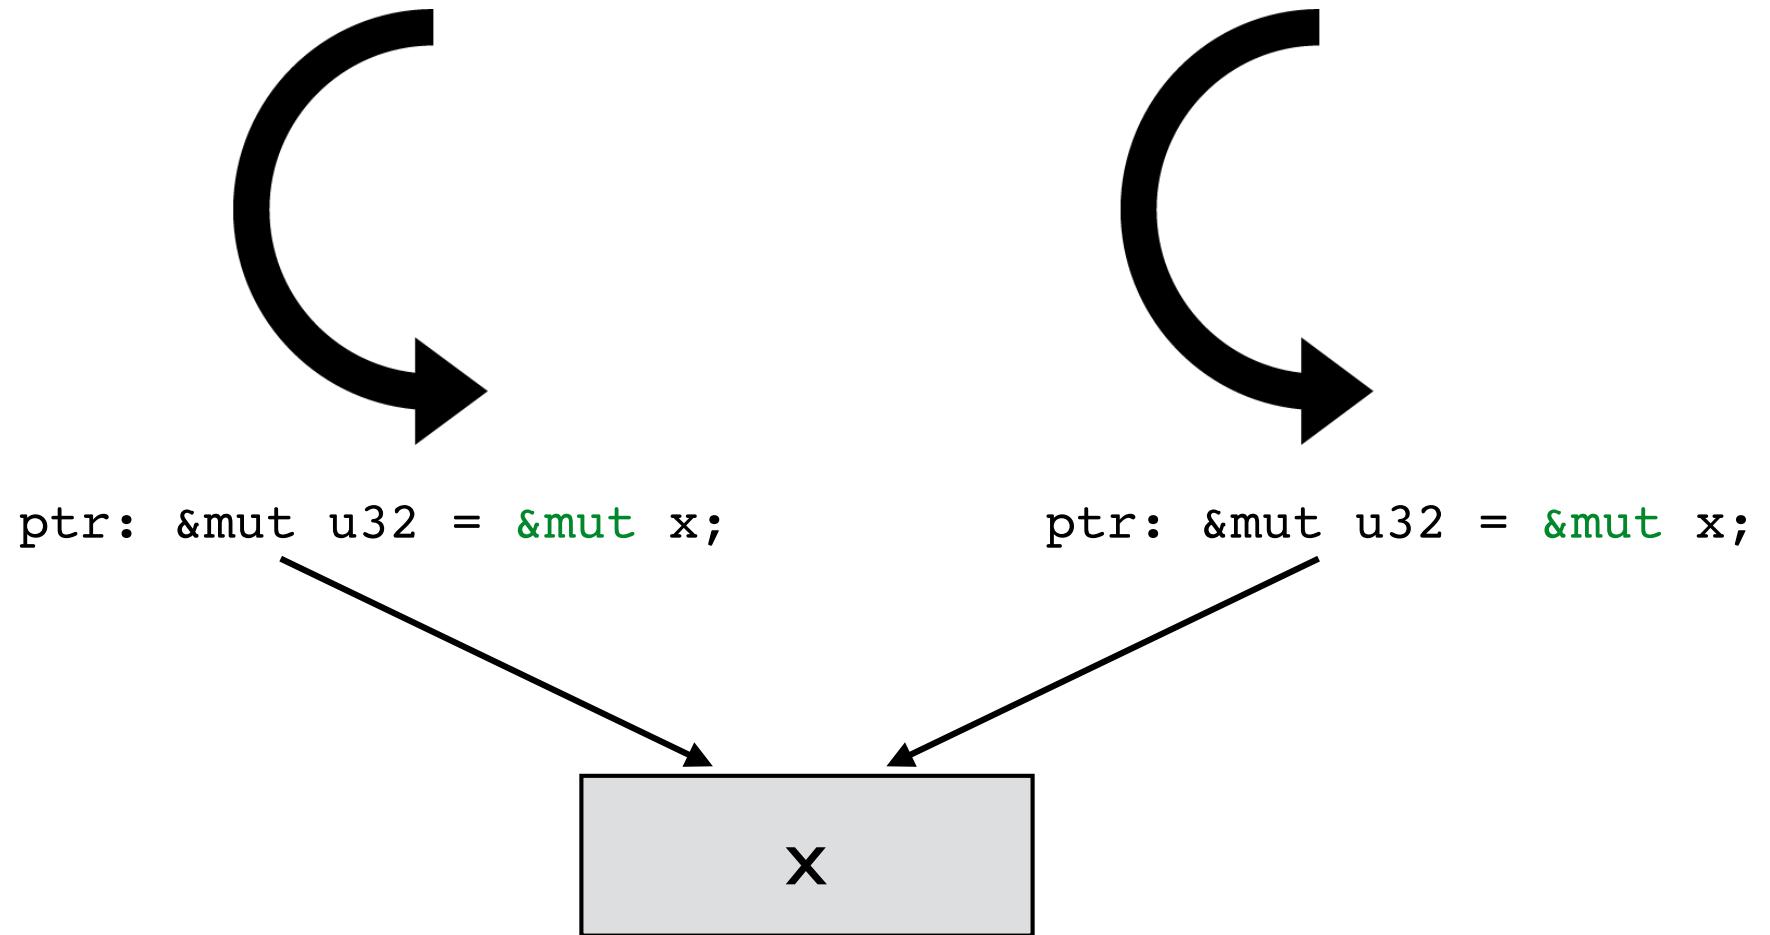
```
let mut x = 5;  
let y = &mut x;  
let z = &x;
```

No

```
let mut x = 5;  
let y = &mut x;  
let z = &mut x;
```

No

Why: Thread Safety



Why: Memory Safety

```
union NumOrPointer {  
    uint32_t Num;  
    uint32_t* Pointer;  
};
```

```
union NumOrPointer* external;  
uint32_t* numptr = &external->Num;  
*numptr = 0xdeadbeef;  
*external->Pointer = 12345;
```

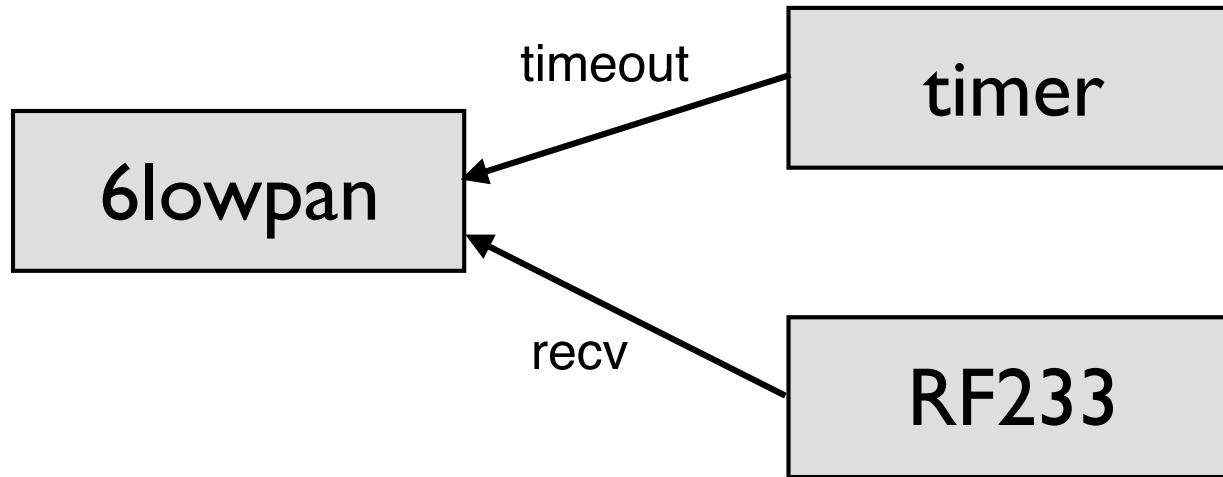
```
enum NumOrPointer {  
    Num(u32),  
    Pointer(&'static mut u32)  
}  
  
// n.b. compile error  
let external : &mut NumOrPointer;  
match external {  
    &mut Pointer(ref mut internal) => {  
        // This would violate safety and  
        // write to memory at 0xdeadbeef  
        *external = Num(0xdeadbeef);  
        *internal = 12345;  
    },  
    ...  
}
```

C

Rust

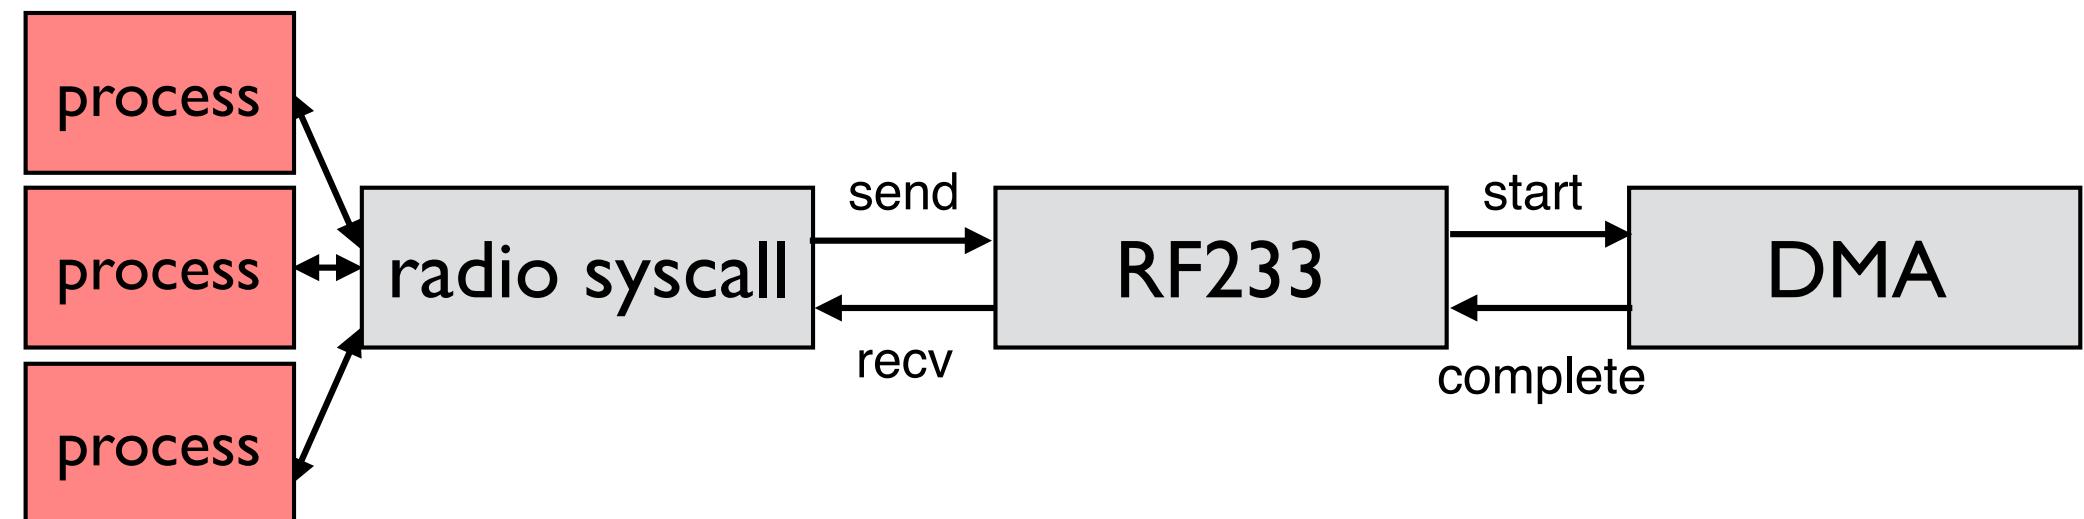
Problem I: Events

- Often want to register multiple event callbacks on a single structure
 - ▶ E.g., networking stack has packet reception and timers
- Each callback needs a mutable reference



Problem 2: System Calls

- Decompose abstractions into kernel components (Rust APIs) with an optional system call interface
- Both upper and lower layers need to be able to call component: two mutable references



Strawman: Change Language

- Problem is that Rust is conflating memory safety and thread safety
- Introduce execution contexts into the language
- Allow multiple mutable references within a single execution context

Good: Solves thread safety.

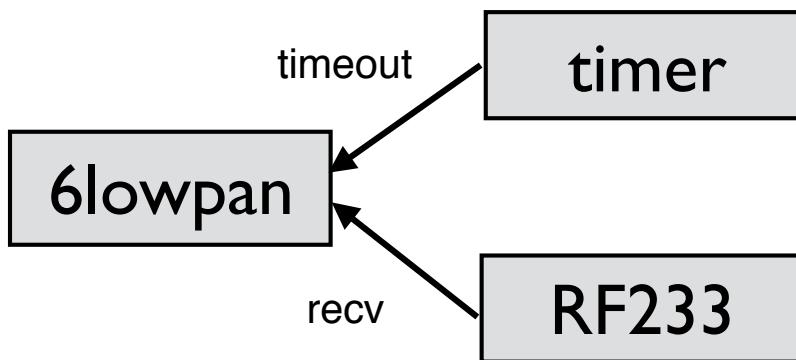
Bad: Doesn't solve memory safety.

```
// Does not compile: Cannot borrow across contexts
let x = 0;
spawn(|| { println!("In thread: {}", x); });

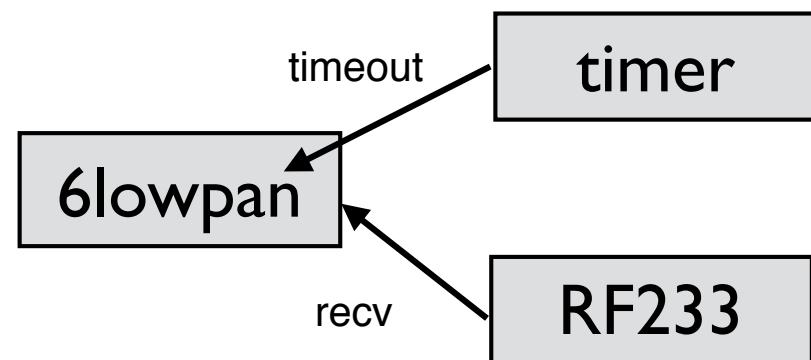
// Does compile: Can take sole ownership with move
let x = 0;
spawn(move || { println!("In thread: {}", x); });
```

Events: Insight

- If we can ensure memory outlives reference, then multiple mutable references can be safe
- Rule: if there is a reference to memory block M, there cannot be any references *inside* M



Safe

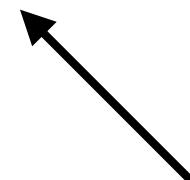


Unsafe

First Step: Cells

- Rust type: allows get/set with immutable references
- Good: multiple references can set values
- Bad: requires memory copies

```
pub struct Spi {  
    registers: *mut SpiRegisters,  
    client: Cell<Option<&'static SpiMasterClient>>,  
    dma_read: Cell<Option<&'static DMAChannel>>,  
    dma_write: Cell<Option<&'static DMAChannel>>,  
    transfers_in_progress: Cell<u8>,  
    dma_length: Cell<usize>,  
}
```

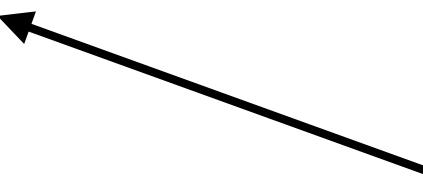


Fields accessed only a few times per operation (async DMA)

First Step: Cells

- Rust type: allows get/set with immutable references
- Good: multiple references can set values
- Bad: requires memory copies

```
pub struct DMAChannel {  
    registers: *mut DMARegisters,  
    nvic: nvic::NvicIdx,  
    pub client: Option<&'static mut DMAClient>,  
    enabled: Cell<bool>,  
    buffer: Cell<'static, [u8]>,  
}
```



Copying a whole buffer!

TakeCell

- Rather than copy data in/out, a TakeCell uses a closure
 - ▶ Pass code in
 - ▶ Thread safety: until first closure completes, other invocations see TakeCell as empty (None)
- Holds a reference to a type

```
struct App { /* many vars */ }  
app: TakeCell<App>,  
  
self.app.map(|app| {  
    // code can read/write  
    // app's variables  
});
```

```
pub struct DMAChannel {  
    registers: *mut DMARegisters,  
    nvic: nvic::NvicIdx,  
    pub client: Option<&'static mut DMAClient>,  
    enabled: Cell<bool>,  
    buffer: TakeCell<'static, [u8]>,  
}
```

Example Code and ASM

```
struct App {  
    count: u32,  
    tx_callback: Callback,  
    rx_callback: Callback,  
    app_read: Option<AppSlice<Shared, u8>>,  
    app_write: Option<AppSlice<Shared, u8>>,  
}  
  
pub struct Driver {  
    busy: Cell<bool>,  
    app: TakeCell<App>,  
}  
driver.app.map(|app| {  
    app.count = app.count + 1  
});
```

/* Load App address into r1, replace with null */
ldr r1, [r0, 0]
movs r2, 0
str r2, [r0, 0]
/* If TakeCell is empty (null) return */
cmp r1, 0
it eq
bx lr
/* Non-null: increment count */
ldr r2, [r1, 0]
add r2, r2, 1
str r2, [r1, 0]
/* Store App back to TakeCell */
str r1, [r0, 0]
bx lr

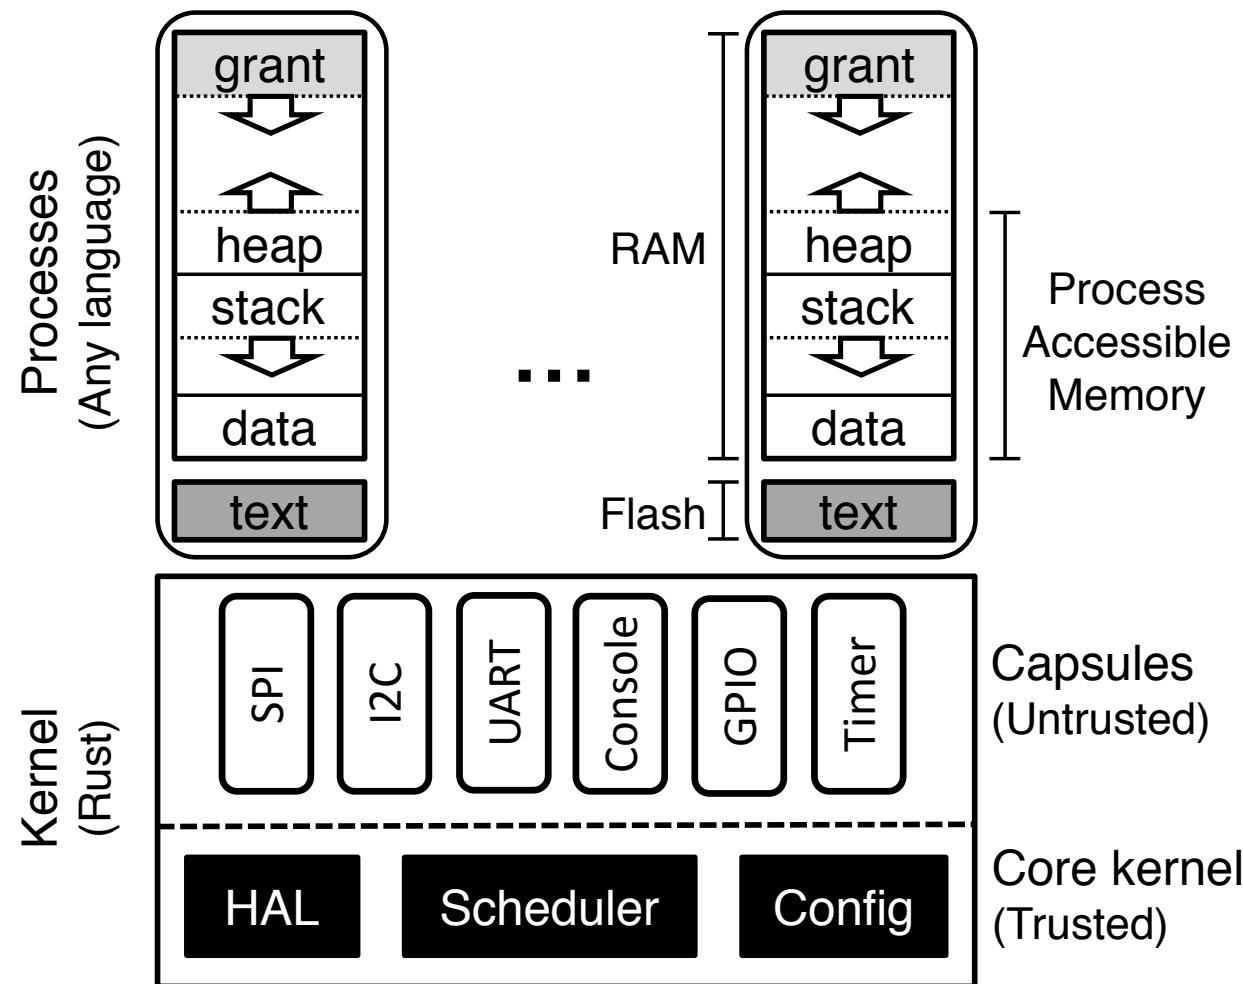
Resource Model

- Kernel is a single thread of control
 - ▶ All I/O operations are non-blocking
 - ▶ Interrupts enqueue events
 - ▶ System calls enqueue events
- Kernel has no heap
 - ▶ Prevent resource exhaustion from shared pool

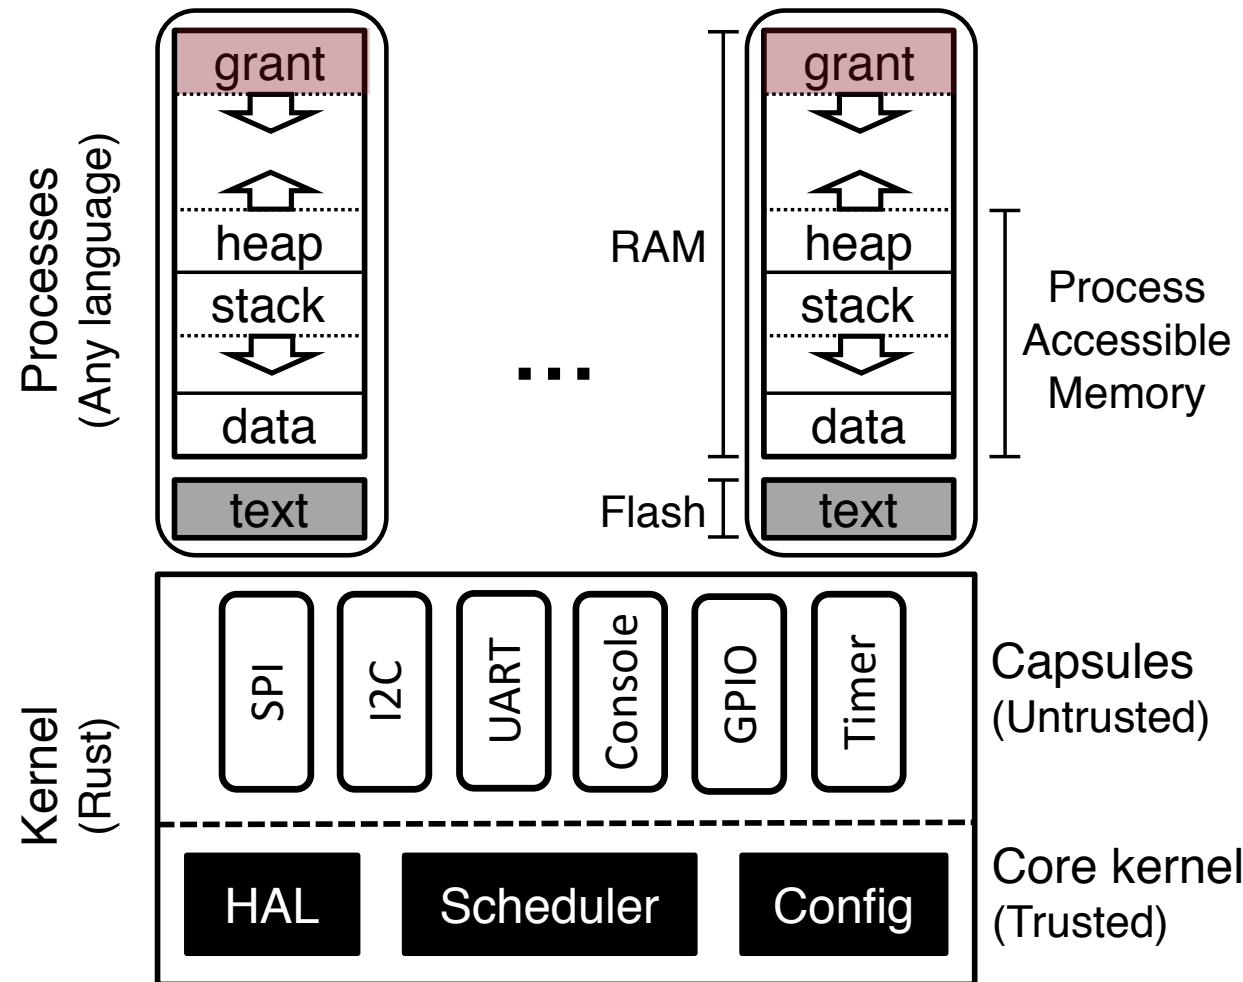
Problem 2: System Calls

- System calls need to dynamically allocate memory
 - ▶ Create a timer, kernel needs to keep timer's state
 - ▶ Enqueue a packet to send, kernel needs reference to packet
- Kernel can't dynamically allocate memory!
 - ▶ Otherwise a process can exhaust kernel memory
 - ▶ Fragmentation
 - ▶ Cleaning up after process failures

System Call Insight



System Call Insight

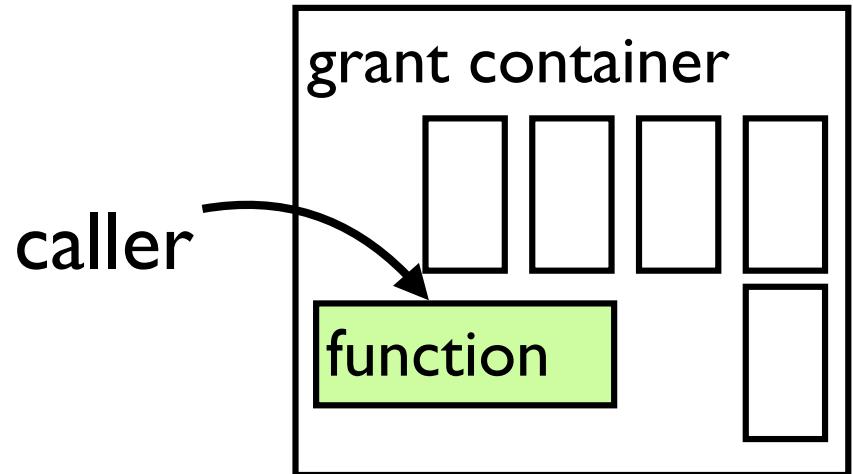


- Processes given block of memory
- Dynamically allocated when process loaded
- Kernel can allocate memory from process

Memory Grants

- Each process has a growable container of *grant memory*
- Kernel can allocate objects from the grant block
- References to objects cannot escape the block
 - ▶ Process failure/crash does not lead to dangling pointers
- Users pass a function to the container with `enter`

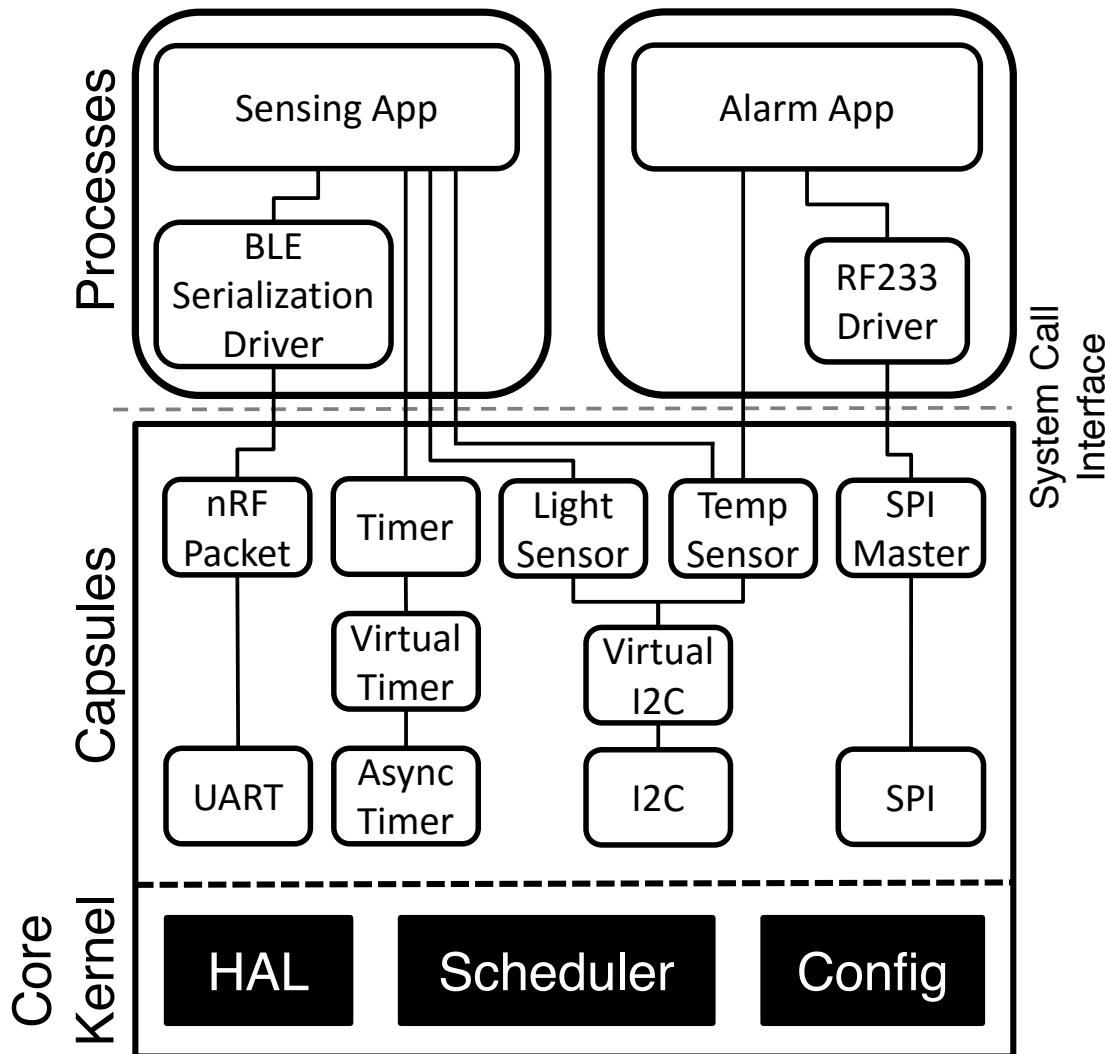
```
self.apps.enter(appid, |app, _| {
    app.read_buffer = Some(slice);
    app.read_idx = 0;
}
).unwrap_or(-1)
```



Outline

- Energy rules everything
- Kernel design goals
- Tock operating system design
 - ▶ Architecture: Rust, capsules, and processes
 - ▶ Rust challenges
 - ▶ Resource model: memory and CPU
- Some early observations

Drivers as Processes



Processes allow easy incorporation of existing library code and drivers:
BLE serialization, 802.15.4 radio, etc.

Truly Non-Blocking

- Prior kernels claimed to be non-blocking
 - ▶ TinyOS
 - ▶ Contiki
 - ▶ SOS
- Used blocking calls for very short operations
 - ▶ E.g., write a single byte over the SPI
- CortexMs are 8x faster: the cycles for these blocking operations is now significant
 - ▶ Every state transition in an FSM is non-blocking...

RF233

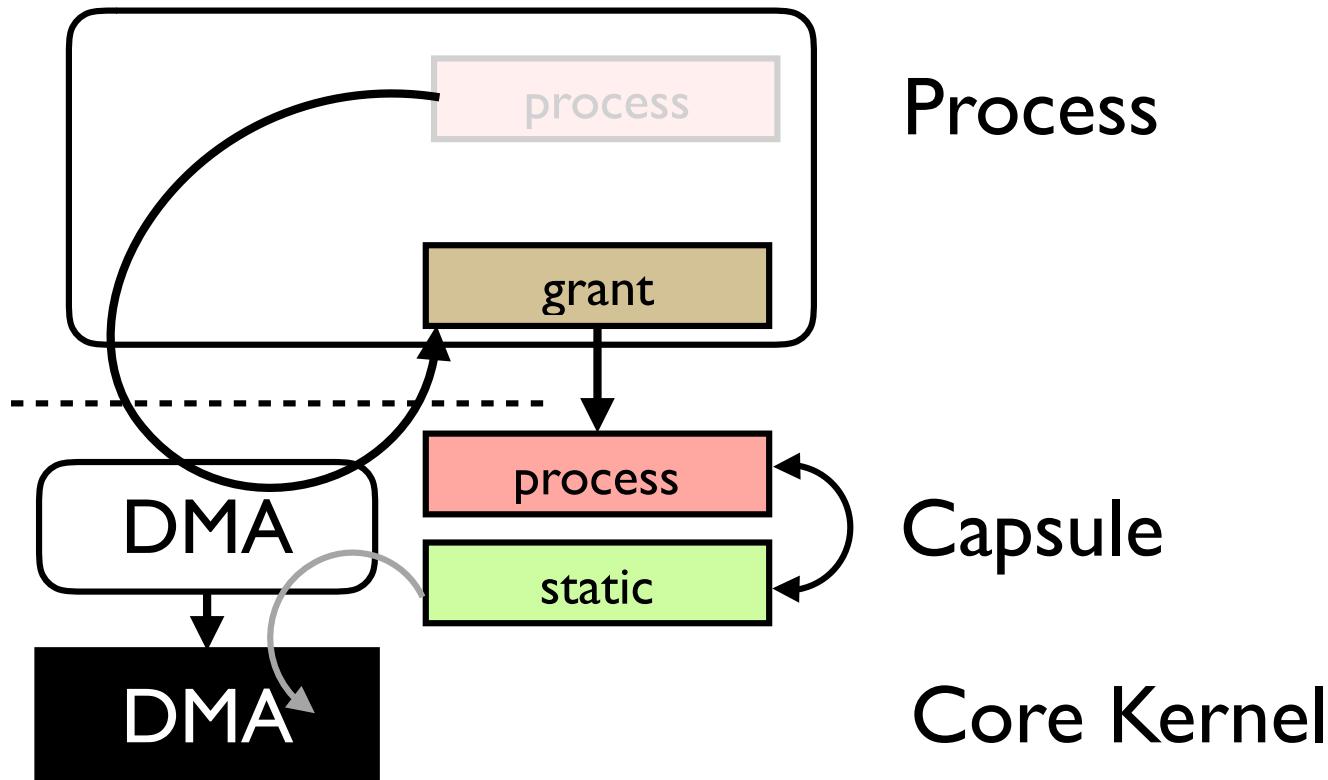
- 802.15.4 radio (used in imix and Nest Protect)
- On an interrupt, have to perform an SPI transaction to read what the source(s) of the interrupt are
 - ▶ You could be in the midst of an existing SPI transaction, such as reading the interrupt state
- Client calls can start complex series of SPI transactions
 - ▶ Sending a packet
 - ▶ Reconfiguring the radio
- End result: 64 states, plus 4 pending flags: transmitting, receiving, configuration, and interrupt

Lifetimes Force Discipline

- A Rust *lifetime* is how long a memory object exists
- If a reference R points to an object O , you must prove that $O_{\text{lifetime}} > R_{\text{lifetime}}$: no dangling pointer
- Lifetime is inherently tied to basic blocks and the stack
- Hardware works differently...
- Question: what's the required lifetime of a buffer whose address you put into a DMA register?

DMA Engine

```
pub fn do_xfer(&self,  
               pid: DMAPeripheral,  
               buf: &'static mut [u8],  
               len: usize) {
```



Tock Operating System

- Safe, multi-tasking operating system for memory-constrained devices
- Core kernel written in Rust, a safe systems language
 - ▶ Small amount of trusted code (can do unsafe things)
 - Rust bindings for memory-mapped I/O
 - Core scheduler, context switches
- Core kernel can be extended with *capsules*
 - ▶ Safe, written in Rust
 - ▶ Run inside kernel
- Processes can be written in any language (asm, C)
 - ▶ Leverage Cortex-M memory protection unit (MPU)
 - ▶ User-level, traps to kernel with system calls

Previous Systems

| | No Resource Exhaustion | Memory Isolation | Memory Efficiency | Concurrency | Loadable Applications |
|-------------|------------------------|------------------|-------------------|-------------|-----------------------|
| Arduino | | | ✓ | | |
| TinyOS | | | ✓ | ✓ | |
| SOS | | | ✓ | ✓ | ✓ |
| TOSTThreads | ✓ | | | ✓ | ✓ |
| Contiki | | | ✓ | ✓ | ✓ |
| FreeRTOS | | | | ✓ | |
| RiotOS | | | ✓ | | |

Previous Systems

| | No Resource Exhaustion | Memory Isolation | Memory Efficiency | Concurrency | Loadable Applications |
|-------------|------------------------|------------------|-------------------|-------------|-----------------------|
| Arduino | | | ✓ | | |
| TinyOS | | | ✓ | ✓ | |
| SOS | | | ✓ | ✓ | ✓ |
| TOSTThreads | ✓ | | | ✓ | ✓ |
| Contiki | | | ✓ | ✓ | ✓ |
| FreeRTOS | | | | ✓ | |
| RiotOS | | | ✓ | | |
| Tock | ✓ | ✓ | ✓ | ✓ | ✓ |

Thanks!

<https://www.tockos.org/>

